

Defining Effect Methods for Other Models

John Fox and Sanford Weisberg

August 24, 2018

The **effects** package in R is designed primarily to draw graphs that visualize a fitted response surface of a fitted model in problems with a linear predictor. Many modeling paradigms that can be fit with base R or contributed packages fit into this framework, including methods for linear, multivariate linear, and generalized linear models fit by the standard **lm** and **glm** functions and by the **svyglm** function in the **survey** package (Lumley, 2004); linear models fit by generalized least squares using the **gls** function in the **nlme** package (Pinheiro et al., 2016); multinomial regression models fit by **multinom** in the **nnet** package (Venables and Ripley, 2002); ordinal regression models using **polr** from the **MASS** package (Venables and Ripley, 2002) and **c1m** and **c1m2** from the **ordinal** package (Christensen, 2015); linear and generalized linear mixed models using the **lme** function in the **nlme** package (Pinheiro et al., 2016) and the **lmer** and **glmer** functions in the **lme4** package (Bates et al., 2015); and latent class models fit by **poLCA** in the **poLCA** package (Linzer and Lewis, 2011). This is hardly an exhaustive list of fitting methods that are based on a linear predictor, and we have been asked from time to time to write functions to use **effects** with this other fitting methods. The mechanism for this is fairly simple. This vignette assumes you are familiar with R's S3 methods.

The default **Effect.default** may work with some modeling functions, as would objects of the class **gls** that we describe below in Section 1, but as illustrated in later sections you may need to modify some of the arguments that are sent to **Effect.default**.

The **effect** package has five functions that create the information needed for drawing effects plots, **Effect**, **allEffects**, **effect** and **predictorEffect** and **predictorEffects**. To add new modeling to the package only a new **Effect** needs to be written; the package will take care of all the other functions.

1 Using effects with Other Modeling Methods, with Generalized Least Squares in the nlme package as an Example

Applying **effects** to other than **lm** and **glm** objects *may* require writing an method for the **Effect** generic function for that type of model object. For

example, the `gls` function in the `nlme` package (Pinheiro et al., 2018) fits linear models via generalized least squares. A call to `gls` creates an object of class `gls`. The following method of `Effect` for `gls` objects finds the information needed to draw effects plots from `gls` objects:

```
Effect.gls <- function(focal.predictors, mod, ...){
  args <- list(
    type = "glm",
    call = mod$call,
    formula = formula(mod),
    family = NULL,
    coefficients = coef(mod),
    vcov = as.matrix(vcov(mod)),
    method=NULL)
  Effect.default(focal.predictors, mod, ..., sources=args)
}
```

This function simply harvests the needed information into a variable `sources`, and then passes to the result to the default `Effect` method. The three required arguments to the method are: `focal.predictors` and `mod` that match the first two arguments of `Effect.default`, and `...` that matches all other arguments. The list `sources` has up to 6 named elements and is created in the function:

type The `effects` package has three basic modeling functions: `type = "glm"`, the default, is used for functions with a univariate response and a linear predictor and possibly a link function. This class includes linear models, generalized linear models, robust regression, generalized least squares fitting, linear and generalized linear mixed effects models, and many others. The `type = "polr"` is used for ordinal regression models, as in the `polr` function in the `MASS` package, and similar methods described below in Section 6. The `type = "multinom"` for multinomial log-linear models as fit by the `multinom` function in `nnet`, and to polytomous latent class models created with the `poLCA` function in the `poLCA` package.

call The `Effect.default` method uses the `call` to harvest additional arguments that it needs. For `type="glm"`, these arguments are `formula`, `data`, `contrasts`, `subset`, `family`, and `offset`, although only the `formula` argument is required. The default is `mod$call` for S3 objects and `mod@call` for S4 objects.

formula In most cases the formula for the linear predictor is returned by `formula(mod)`, the default, but if this is not the case the value of this argument should be the value of the formula for fixed effects.

family The default is `family=NULL`. This argument is required for GLM-like models that include a `family` that specifies both an error distribution and a link function only if `family=family(mod)` is not appropriate. See the `betareg` example in Section 5 below for an example that includes a user-selected link function, but a fixed error distribution.

coefficients In many cases the (fixed-effect) coefficient estimates are returned by `coef(mod)`, the default, but if this is not the case then the value of this argument should be the estimates of the coefficients in the linear predictor. The functions in the **effects** package do not use estimates of random effects.

vcov In many cases the estimated covariance matrix of the (fixed-effect) coefficient estimates is returned by `vcov(mod)`, the default, but if this is not the case then the value of this argument should be the estimated covariance matrix of the (fixed-effect) coefficient estimates in the linear predictor.

method This argument is used only for methods that use effects graphics based on the `polr` function, where the argument **method** is the name of a link function; see `help(polr)` for a list of the accepted links, and see Section 6.1 below for an example.

*Since the values of all the arguments in **sources** are default values for the `gls` method, there is no need to have written the `Effect.gls` method, as the default method would work.*

```
library(effects)
```

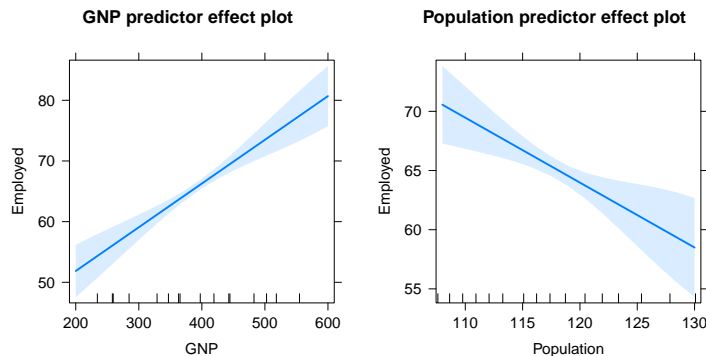
Loading required package: carData

```
lattice theme set by effectsTheme()
See ?effectsTheme for details.
```

```
require(nlme)
```

Loading required package: nlme

```
g <- gls(Employed ~ GNP + Population,
        correlation=corAR1(form= ~ Year), data=longley)
plot(predictorEffects(g))
```



2 Mixed Effects with lme (nlme package)

The `lme` function in the `nlme` package (Pinheiro et al., 2018) fits linear mixed models. The required function for fitted objects from this function is included in the `effects` package. It is given by

```
print(Effect.lme)

function(focal.predictors, mod, ...){
  args <- list(
    call = mod$call,
    formula = mod$call$fixed,
    coefficients = mod$coefficients$fixed,
    vcov = mod$varFixed)
  Effect.default(focal.predictors, mod, ..., sources=args)
}
<bytecode: 0x7f803aec0320>
<environment: namespace:effects>
```

The `formula`, `coefficients` and `vcov` arguments are set to non-default values. The other arguments are automatically set to default values.

```
data(Orthodont, package="nlme")
m1 <- nlme::lme(distance ~ age + Sex, data=Orthodont,
               random= ~ 1 | Subject)
as.data.frame(Effect("age", m1))

   age      fit      se   lower   upper
1  8.0 22.04259 0.4172841 21.21520 22.86999
2  9.5 23.03287 0.3853671 22.26876 23.79698
3 11.0 24.02315 0.3741236 23.28133 24.76497
4 12.0 24.68333 0.3791619 23.93153 25.43514
5 14.0 26.00370 0.4172841 25.17631 26.83110
```

3 Mixed Effects with the lmer (lme4 package)

The `lme4` package (Bates et al., 2015) fits linear and generalized linear mixed effects models with the `lmer` and `glmer` functions, respectively. The same `Effect` function can be used for `lmer` and `glmer` models.

The following method is a little more complicated because it contains an additional argument `KR` to determine if the Kenward-Roger coefficient covariance matrix is to be used to compute effect standard errors. The default is `FALSE` because the computation is very slow. If `KR = TRUE`, the function also checks if the `pbkrtest` package is present.

```
print(Effect.merMod)
```

```

function(focal.predictors, mod, ..., KR=FALSE){
  if (KR && !requireNamespace("pbkrtest", quietly=TRUE)){
    KR <- FALSE
    warning("pbkrtest is not available, KR set to FALSE")
  }
  fam <- family(mod)
  args <- list(
    call = mod@call,
    coefficients = lme4::fixef(mod),
    family=fam,
    vcov = if (fam == "gaussian" && fam$link == "identity" && KR)
      as.matrix(pbkrtest::vcovAdj(mod)) else as.matrix(vcov(mod))
  )
  Effect.default(focal.predictors, mod, ..., sources=args)
}
<bytecode: 0x7f803ffc49b0>
<environment: namespace:effects>

```


Because `lmer` is an S4 object, the default for `call` is `mod@call`, and this argument would have been set automatically had we not included it in the above method. The fixed-effect estimates for an object created by a call to `lmer` or `glmer` are not returned by `coef(mod)`, so the value of `coefficients` is the value returned by `lme4::fixef(mod)`. The `vcov` estimate contains its estimated variance covariance matrix of the fixed effects. The Kenward-Roger method is used to estimate the covariance matrix for linear models if the additional argument `KR=TRUE`. The default is `KR=FALSE` because The Kenward-Roger estimate requires a long computation; see `help(Effect)`.

The `formula` for a mixed-effects model in the `lme4` package specifies linear predictors for both the mean function and the variance functions, specified by, for example `(1 + age | Subject)`. The `effects` code will automatically remove any terms like these in any formula, as the effects package only displays the mean function.

```

fm2 <- lme4::lmer(distance ~ age + Sex + (1 |Subject), data
                  = Orthodont)
plot(allEffects(fm2))

```



```
data(cbpp, package="lme4")
gm1 <- lme4::glmer(cbind(incidence, size - incidence) ~ period + (1 | herd),
  data = cbpp, family = binomial)
as.data.frame(predictorEffect("period", gm1))
```

	period	fit	se	lower	upper
1	1	0.19807921	0.2312141	0.13569522	0.2798569
2	2	0.08391784	0.3073939	0.04775453	0.1433443
3	3	0.07401714	0.3270802	0.04040242	0.1317591
4	4	0.04842565	0.4251645	0.02163871	0.1048199

4 Robust Linear Mixed Models (robustlmm package)

The `rlmer` function in the `robustlmm` package (Koller, 2016) fits linear mixed models with a robust estimation method. As `rlmer` closely parallels the `lmer` function, an object created by `rlmer` is easily used with `effects`:

```
print(Effect.rlmerMod)

function(focal.predictors, mod, ...){
  args <- list(
    coefficients = lme4::fixef(mod),
    family=family(mod))
  Effect.default(focal.predictors, mod, ..., sources=args)
}
<bytecode: 0x7f80422fa1c8>
<environment: namespace:effects>

require(lme4)
fm3 <- robustlmm::rlmer(distance ~ age * Sex + (1 |Subject),
  data = Orthodont)
plot(predictorEffects(fm3))
```

5 Beta Regression

The `betareg` function in the `betareg` package (Grün et al., 2012) fits regressions with a link function but with Beta distributed errors.

```
print(Effect.betareg)

function(focal.predictors, mod, ...){
  coef <- mod$coefficients$mean
  vco <- vcov(mod)[1:length(coef), 1:length(coef)]
  # betareg uses beta errors with mean link given in mod$link$mean.
  # Construct a family based on the binomial() family
```

```

    fam <- binomial(link=mod$link$mean)
  # adjust the varince function to account for beta variance
  fam$variance <- function(mu){
    f0 <- function(mu, eta) (1-mu)*mu/(1+eta)
    do.call("f0", list(mu, mod$coefficient$precision))}
  # adjust initialize
  fam$initialize <- expression({mustart <- y})
  args <- list(
    call = mod$call,
    formula = formula(mod),
    family=fam,
    coefficients = coef,
    vcov = vco)
  Effect.default(focal.predictors, mod, ..., sources=args)
}
<bytecode: 0x7f803f693270>
<environment: namespace:effects>

```

Beta regression has a response $y \in [0, 1]$, with the connection between the mean μ of the Beta and a set for predictors \mathbf{x} through a link function $\mathbf{x}'\boldsymbol{\beta} = g(\mu)$. The variance function for the beta is $\text{var}(y) = \mu(1 - \mu)/(1 + \phi)$, for a precision parameter ϕ estimated by **betareg**.

The call to **betareg** does not have a family argument, although it does have a link stored in `mod$link$mean`. For use with **Effect.default**, the method above creates a family from the binomial family generator. It then adjusts this family by changing from binomial variance to the variance for the beta distribution. Since the **glm** function expects a variance that is a function of only one parameter, we fix the value of the precision ϕ at its estimator from the **betareg** fit, as shown in the method. We need to replace the **initialize** method in the family to one appropriate for $y \in [0, 1]$.

```

library(betareg)
require(lme4)
data("GasolineYield", package = "betareg")
gy_logit <- betareg(yield ~ batch + temp, data = GasolineYield)
summary(gy_logit)

```

Call:

```
betareg(formula = yield ~ batch + temp, data = GasolineYield)
```

Standardized weighted residuals 2:

Min	1Q	Median	3Q	Max
-2.8750	-0.8149	0.1601	0.8384	2.0483

Coefficients (mean model with logit link):

Estimate	Std. Error	z value	Pr(> z)
----------	------------	---------	----------

(Intercept)	-6.1595710	0.1823247	-33.784	< 2e-16
batch1	1.7277289	0.1012294	17.067	< 2e-16
batch2	1.3225969	0.1179020	11.218	< 2e-16
batch3	1.5723099	0.1161045	13.542	< 2e-16
batch4	1.0597141	0.1023598	10.353	< 2e-16
batch5	1.1337518	0.1035232	10.952	< 2e-16
batch6	1.0401618	0.1060365	9.809	< 2e-16
batch7	0.5436922	0.1091275	4.982	0.000000629
batch8	0.4959007	0.1089257	4.553	0.000005297
batch9	0.3857930	0.1185933	3.253	0.00114
temp	0.0109669	0.0004126	26.577	< 2e-16

Phi coefficients (precision model with identity link):

	Estimate	Std. Error	z value	Pr(> z)
(phi)	440.3	110.0	4.002	0.0000629


Type of estimator: ML (maximum likelihood)

Log-likelihood: 84.8 on 12 Df

Pseudo R-squared: 0.9617

Number of iterations: 51 (BFGS) + 3 (Fisher scoring)

`plot(predictorEffects(gy_logit))`



6 Ordinal Models (ordinal package)

Proportional odds logit and probit regression models fit with the `polr` function in the `MASS` package (Venables and Ripley, 2002) are supported in the `effects` package. The `ordinal` package, (Christensen, 2015) contains three functions that are very similar to `polr`. The `clm` and `clm2` functions allow more link functions and a number of other generalizations. The `clmm` function allows including random effects.


6.1 clm

```
print(Effect.clm)

function(focal.predictors, mod, ...){
  if (requireNamespace("MASS", quietly=TRUE)){
    polr <- MASS::polr} else stop("MASS package is required")
  polr.methods <- c("logistic", "probit", "loglog",
                    "cloglog", "cauchit")
  method <- mod$link
  if(method == "logit") method <- "logistic"
  if(!(method %in% polr.methods))
    stop("'link' must be a 'method' supported by polr; see help(polr)")
  if(mod$threshold != "flexible")
    stop("Effects only supports the 'flexible' threshold")
  if(is.null(mod$Hessian)){
    message("\nRe-fitting to get Hessian\n")
    mod <- update(mod, Hess=TRUE)}
  numTheta <- length(mod$Theta)
  numBeta <- length(mod$beta)
  or <- c( (numTheta+1):(numTheta + numBeta), 1:(numTheta))
  args <- list(
    type = "polr",
    coefficients = mod$beta,
    method=method,
    vcov = as.matrix(vcov(mod)[or, or]))
  Effect.default(focal.predictors, mod, ..., sources=args)
}
<bytecode: 0x7f8027d56498>
<environment: namespace:effects>
```

This method first checks that the MASS package is available. Since the `clm` function allows suppressing the computation of the Hessian, the function checks and computes it if needed to get the estimated covariance matrix. The `clm` function orders the parameters in the order (threshold parameters, linear predictor parameters), so the next few lines identify the elements of `vcov` that are needed by `Effects`. Since the `polr` function does not allow thresholds other than `flexible`, we don't allow them either. The `polr` argument `method` is equivalent to the `clm` argument `link`, except that the `clm` link "logit" is equivalent to the `polr` method "logit" "logistic".

```
require(ordinal)
require(MASS)
mod.wvs1 <- clm(poverty ~ gender + religion + degree + country*poly(age,3),
               data=WVS)
plot(Effect(c("country", "age"), mod.wvs1),
     lines=list(multiline=TRUE), layout=c(2, 2))
```



6.2 clm2

Although the fitted models are similar, syntax for `clm2` is not the same as `clm`, so a separate method is required.

```
print(Effect.clm2)


function(focal.predictors, mod, ...){
  if (requireNamespace("MASS", quietly=TRUE)){
    polr <- MASS::polr}
  polr.methods <- c("logistic", "probit", "loglog",
                    "cloglog", "cauchit")
  method <- mod$link
  if(!(method %in% polr.methods))
    stop("'link' must be a 'method' supported by polr; see help(polr)")
  if(is.null(mod$Hessian)){
    message("\nRe-fitting to get Hessian\n")
    mod <- update(mod, Hess=TRUE)}
  if(mod$threshold != "flexible")
    stop("Effects only supports the flexible threshold")
  numTheta <- length(mod$Theta)
  numBeta <- length(mod$beta)
```

```

or <- c( (numTheta+1):(numTheta + numBeta), 1:(numTheta))
args <- list(
  type = "polr",
  formula = mod$call$location,
  coefficients = mod$beta,
  method=method,
  vcov = as.matrix(vcov(mod)[or, or]))
Effect.default(focal.predictors, mod, ..., sources=args)
}
<bytecode: 0x7f80469308e8>
<environment: namespace:effects>

v2 <- clm2(poverty ~ gender + religion + degree + country*poly(age,3),data=WVS)
plot(emod2 <- Effect(c("country", "age"), v2))

```



6.3 clmm

This function allows for random effects in an ordinal model.

```

print(Effect.clmm)

function(focal.predictors, mod, ...){
  if (requireNamespace("MASS", quietly=TRUE)){
    polr <- MASS::polr}
  else stop("The MASS package must be installed")
  polr.methods <- c("logistic", "probit", "loglog",
                    "cloglog", "cauchit")
  method <- mod$link

```

```

if(method == "logit") method <- "logistic"
if(!(method %in% polr.methods))
  stop("'link' must be a 'method' supported by polr; see help(polr)")
if(is.null(mod$Hessian)){
  message("\nRe-fitting to get Hessian\n")
  mod <- update(mod, Hess=TRUE)}
if(mod$threshold != "flexible")
  stop("Only threshold='flexible' supported by Effects")
numTheta <- length(mod$Theta)
numBeta <- length(mod$beta)
or <- c( (numTheta+1):(numTheta + numBeta), 1:(numTheta))
Vcov <- as.matrix(vcov(mod)[or, or])
args <- list(
  type = "polr",
  formula = formula(mod),
  coefficients = mod$beta,
  method=method,
  vcov = as.matrix(Vcov))
Effect.default(focal.predictors, mod, ..., sources=args)
}
<bytecode: 0x7f8027b13930>
<environment: namespace:effects>

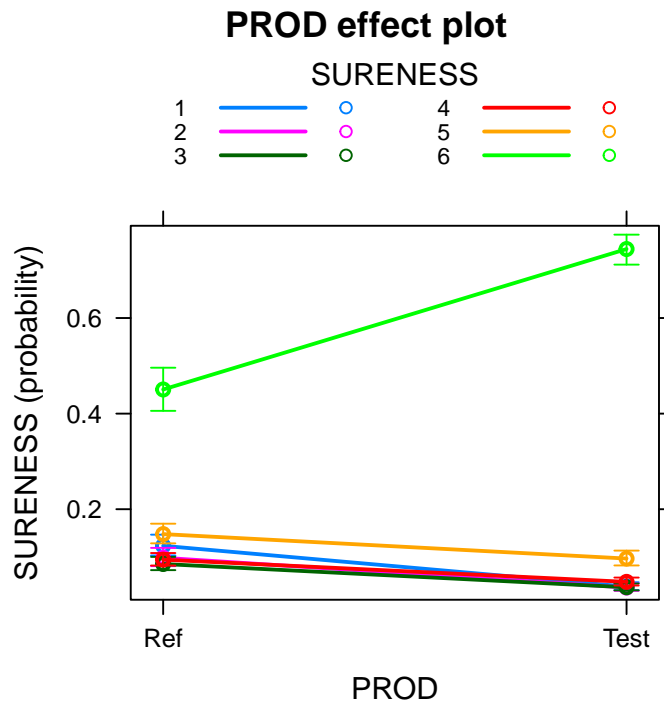
```

The first few lines of the method check for the presence of the MASS package that is needed to use polr, makes sure the link used is supported by polr, and requires that the argument `threshold` has its default value. The `polr` and `clmm` functions store the fixed effects estimates of regression and threshold coefficients in different orders, so the next few lines rearrange the variance matrix to match the order that polr uses.

```

require(ordinal)
require(MASS)
mm1 <- clmm(SURENESS ~ PROD + (1|RESP) + (1|RESP:PROD),
  data = soup, link = "logit", threshold = "flexible")
plot(Effect("PROD", mm1), lines=list(multiline=TRUE))

```



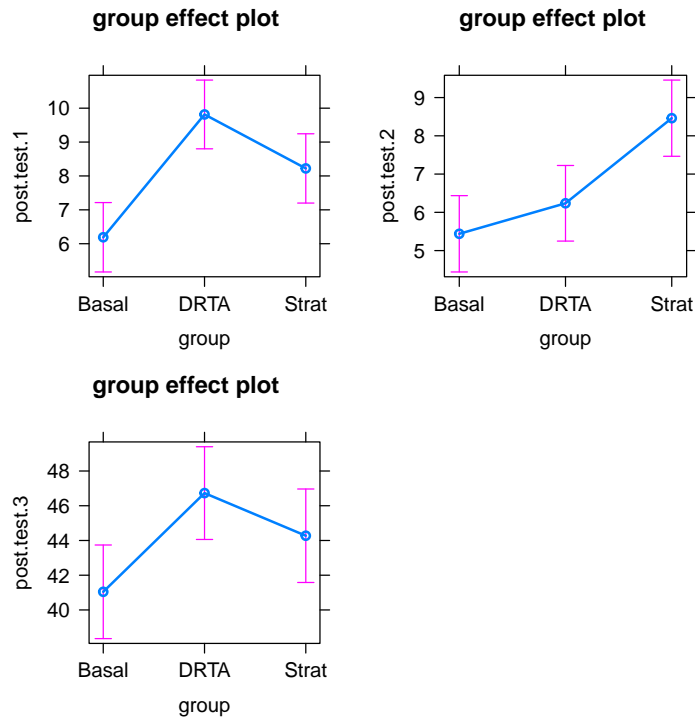
6.4 Others

The `poLCA` function in the `poLCA` package (Linzer and Lewis, 2011) fits polytomous variable latent class models, which uses the multinomial effects plots.

The `svyglm` function in the `survey` package (Lumley, 2004, 2016) fits generalized linear models using survey weights.

The `lm` function can also be used to create a multivariate linear model. The `Effect.mlm` function, with slightly different syntax, will draw effects plots for these models, with separate plots of each response.

```
data(Baumann, package="carData")
b1 <- lm(cbind(post.test.1, post.test.2, post.test.3) ~ group +
         pretest.1 + pretest.2, data = Baumann)
plot(Effect("group", b1))
```



References

- Bates, D., M. Mächler, B. Bolker, and S. Walker (2015). Fitting linear mixed-effects models using lme4. *Journal of Statistical Software* 67(1), 1–48.
- Christensen, R. H. B. (2015). `ordinal`—Regression Models for Ordinal Data. R package version 2015.6-28.
- Grün, B., I. Kosmidis, and A. Zeileis (2012). Extended beta regression in R: Shaken, stirred, mixed, and partitioned. *Journal of Statistical Software* 48(11), 1–25.
- Koller, M. (2016). `robustlmm`: An R package for robust estimation of linear mixed-effects models. *Journal of Statistical Software* 75(6), 1–24.
- Linzer, D. A. and J. B. Lewis (2011). `poLCA`: An ^o package for polytomous variable latent class analysis. *Journal of Statistical Software* 42(10), 1–29.
- Lumley, T. (2004). Analysis of complex survey samples. *Journal of Statistical Software* 9(1), 1–19. R package version 2.2.
- Lumley, T. (2016). `survey`: analysis of complex survey samples. R package version 3.32.

- Pinheiro, J., D. Bates, S. DebRoy, D. Sarkar, and R Core Team (2016). `nlme`: Linear and Nonlinear Mixed Effects Models. R package version 3.1-127.
- Pinheiro, J., D. Bates, S. DebRoy, D. Sarkar, and R Core Team (2018). *nlme: Linear and Nonlinear Mixed Effects Models*. R package version 3.1-137.
- Venables, W. N. and B. D. Ripley (2002). *Modern Applied Statistics with S* (4th ed.). New York: Springer-Verlag.